

Ask the Sensors: Grounded, Explainable Activity Question Answering from Wearable Signals

CS60055 Ubiquitous Computing | Hackathon Challenge 1

Author: Anubhav (IIT Kharagpur) | **Teaching Team:** sandipc-iitkgp, sayantan-kuila, debjit2001

Ask the Sensors: Grounded, Explainable Activity Question Answering from Wearable Signals

Course: CS60055 Ubiquitous Computing | Hackathon Challenge 1

Institution: Department of Computer Science and Engineering, IIT Kharagpur

Teaching Team Reviewers: sandipc-iitkgp, sayantan-kuila, debjit2001

Abstract

Wearable continuous sensor streams offer rich diagnostic and behavioral insights, yet converting high-frequency, noisy triaxial accelerometer and gyroscope telemetry into trusted, clinical-grade answers remains a fundamental challenge in ubiquitous computing. Conventional human activity recognition (HAR) models produce ungrounded point predictions without temporal extents, duration bounds, or physical justification. In this work, I present **Ask the Sensors**, an end-to-end, multi-tiered sensor question-answering architecture that operates over continuous 25 Hz 6-channel IMU recordings. My system couples a deterministic temporal interval aggregator with a dual recognition backbone: a lightweight 1D convolutional neural network (with INT8 quantization) and an event-driven **Neuromorphic Spiking Neural Network (SNN)** employing delta-spike modulation and Leaky Integrate-and-Fire (LIF) dynamics. To guarantee mathematical faithfulness, Tasks 1–3 (Identification, Temporal Reasoning, and Evidence Grounding) are executed via symbolic temporal algebra over a smoothed activity timeline, yielding zero hallucination on counts, durations, and onset timestamps. Task 4 (Open-World Semantic Reasoning) is resolved via a kinematic physics engine evaluating cadence harmonics, heel-strike impulse attenuation, and gravitational tilt. On the benchmark derived from the in-the-wild ExtraSensory dataset, my pipeline achieves an overall macro-QA accuracy of **91.3%** and grounding acceptance of **86.0%** at IoU ≥ 0.5 . Crucially, my neuromorphic SNN demonstrates an **$16.7\times$ energy reduction** ($1.1\mu\text{J}$ vs $18.4\mu\text{J}$ per query) over dense FP32 execution during sedentary bouts, establishing a Pareto-optimal operating point for resource-constrained edge wearables.

1. Problem Formulation and Clinical Motivation

1.1 The Ubicomp Clinical Dilemma

Consider the canonical scenario: Aparna, a 72-year-old living alone, wears a commercial smartwatch while her son and physiotherapist monitor her functional recovery remotely. Key clinical inquiries naturally arise: - *"Did she take her usual morning walk?"* - *"How much of the afternoon did she spend resting?"* - *"Did she begin running or doing something strenuous at noon?"* - *"Was she using a wheeled or pedal-based mode of movement?"*

While her smartwatch continuously samples 3-axis acceleration ($\mathbf{a}(t) = [a_x, a_y, a_z]^T$) and 3-axis angular velocity ($\boldsymbol{\omega}(t) = [\omega_x, \omega_y, \omega_z]^T$), raw floating-point time series are uninterpretable to clinicians and family members. Standard black-box classifiers assign isolated categorical labels at arbitrary sampling intervals, failing to answer: 1. **Temporal Boundaries:** Exactly when did the bout start and finish? 2. **Cumulative Durations and Bouts:** Did multiple pauses interrupt the activity? 3. **Evidence Grounding:** *On what basis* did the algorithm decide it was walking rather than standing or cycling?

In high-stakes health monitoring, decisions cannot rest on ungrounded model verdicts. The reasoning must trace directly back to physical signal evidence: dominant step frequencies, peak impacts, and sensor modalities.

1.2 Multi-Tiered Task Definition

The system answers natural-language queries across four tiers of increasing difficulty: - **Task 1: Activity Identification.** Categorical recognition and binary verification (e.g., *"What activity is the user performing?"*, *"Is the user running?"*). - **Task 2: Temporal & Quantitative Reasoning.** Cumulative duration calculation, occurrence counting, onset localization, and duration comparisons over multi-hour recordings. - **Task 3: Evidence Grounding.** Directly evaluated on the precision of supporting signal intervals via Intersection-over-Union ($\text{IoU} \geq 0.5$), sensor modality verification, channel attribution, and signal-grounded rationales. - **Task 4: Open-World Semantic Reasoning.** Kinematic deduction of unmodeled, continuous, or out-of-vocabulary behaviors (e.g., pedal-based locomotion, prolonged lying down vs transient pauses) derived from first-principles mechanics.

2. In-the-Wild Wearable Data Challenges (ExtraSensory)


The ExtraSensory dataset was gathered from 60 users during unconstrained daily routines. Unlike pristine laboratory datasets, it introduces substantial real-world telemetry noise: 1. **Sampling Rate Irregularity & Missing Stretches:** Watch sensors experience OS transmission latency, Bluetooth dropouts, and battery saving throttling. 2. **Heavy Sedentary Class Imbalance:** Users spend upwards of 80% of daily life sitting or lying down; dynamic activities like running or bicycling are

sparsely distributed. 3. **Self-Reported Label Noise & Co-occurrences:** Annotations reflect subjective recall, leading to misaligned timestamps and overlapping labels.

2.1 Uniform 25 Hz Stream Ingestion

To enforce a uniform cross-system time base, every incoming sensor stream is resampled onto an exact $\Delta t = 0.04\text{s}$ ($f_s = 25.0\text{Hz}$) grid: $t_k = t_0 + k \cdot \Delta t, \quad k \in \{0, 1, \dots, N-1\}$ Short sensor dropouts ($\leq 3.0\text{s}$) are repaired using piecewise linear interpolation: $x(t_k) = x(t_i) + \frac{x(t_{i+1}) - x(t_i)}{t_{i+1} - t_i} (t_k - t_i)$ Extended dropouts ($> 3.0\text{s}$) are flagged to prevent spurious motion artifact interpolation. A 4th-order low-pass Butterworth filter ($f_{\text{cutoff}} = 11.0\text{Hz}$), safely below the Nyquist limit of 12.5Hz) attenuates high-frequency electronic jitter and watch casing vibrations.

3. System Architecture & Methodology



3.1 Biomechanical Feature Engineering

For each sliding analysis window $W \in \mathbb{R}^{64 \times 6}$ (2.56 seconds at 25 Hz with 50% overlap), I extract 20 kinematic features:

- Time Domain Statistics:** Vector magnitude mean, variance, standard deviation, and peak acceleration: $\|\mathbf{a}(t)\| = \sqrt{a_x^2(t) + a_y^2(t) + a_z^2(t)}$
- Signal Magnitude Area (SMA):** Energy proxy normalized over window length $L = 64$: $\text{SMA}(\text{acc}) = \frac{1}{L} \sum_{t=1}^L \left(|a_x(t)| + |a_y(t)| + |a_z(t)| \right)$
- Dynamic Jerk:** Rate of change of linear acceleration, isolating shock impacts: $j(t) = \frac{\|\mathbf{a}(t)\| - \|\mathbf{a}(t-1)\|}{\Delta t}$, $\text{Var}(j)$
- Dominant Cadence Frequency via FFT:** Power spectral density estimation along dynamic axes over the gait band $[0.3 \text{ Hz}, 5.0 \text{ Hz}]$: $f_{\text{dom}} = \arg\max_{f \in [0.3, 5.0]} |\mathcal{F}\{\mathbf{a}\}| - \mu_{\{\mathbf{a}\}}(f)$
- Static Gravitational Tilt:** Postural pitch and roll angles derived from low-pass gravity alignment: $\theta_{\text{pitch}} = \arcsin\left(\frac{-\bar{a}_x}{g}\right)$, $\phi_{\text{roll}} = \arctan2(\bar{a}_y, \bar{a}_z)$

4. Neuromorphic Spiking Neural Network (SNN) Innovation

4.1 Motivation: Edge Quiescence in Wearables

Standard deep learning architectures execute continuous floating-point Multiply-Accumulate (MAC) operations regardless of user activity. However, in wearable health monitoring, sedentary states (sitting/lying) dominate over 80% of daily time series. Running dense matrix multiplications on quiescent sensor streams wastes substantial battery energy.

4.2 Delta-Modulated Leaky Integrate-and-Fire (LIF) SNN

To resolve this inefficiency, I developed a neuromorphic Spiking Neural Network (SNN) operating directly on IMU signals:

- Temporal Delta Spike Encoding:** Continuous IMU channels are converted into asynchronous event spikes whenever absolute first-order differences exceed threshold $\delta = 0.15$: $S_i^+(t) = \Theta(x_i[t] - x_i[t-1] - \delta)$, $S_i^-(t) = \Theta(-(x_i[t] - x_i[t-1]) - \delta)$ where Θ is the Heaviside step function. This yields a sparse binary tensor $\mathbf{S}(t) \in \{0, 1\}^{64 \times 12}$.
- LIF Membrane Potential Dynamics:** $V_j[t] = \beta V_j[t-1] \cdot (1 - S_j^{\text{out}}[t-1]) + \sum_i W_{ij} S_i[t]$ where $\beta = 0.85$ represents the membrane leak factor. When $V_j[t] \geq V_{\text{th}} = 1.0$, a spike $S_j^{\text{out}}[t] = 1$ is emitted, and the membrane resets.
- Surrogate Gradient Backpropagation:** Since Θ has zero derivative almost everywhere, I employ the Fast Sigmoid surrogate function during training: $\frac{\partial S}{\partial V} = \frac{1}{(1 + \gamma |V - V_{\text{th}}|)^2}$, $\gamma = 10.0$

4.3 Energy Consumption Modeling: SynOps vs MACs

In classical ANNs, each connection incurs a 32-bit floating-point Multiply-Accumulate operation ($E_{\text{MAC}} \approx 4.6 \text{ pJ}$) in 45nm CMOS). In neuromorphic SNNs, non-spiking timesteps incur zero computation; active spikes trigger simple integer synaptic additions ($E_{\text{AC}} \approx 0.1 \text{ pJ}$): $E_{\text{ANN}} = N_{\text{MAC}} \times 4.6 \text{ pJ}$, $E_{\text{SNN}} = N_{\text{SynOps}} \times 0.1 \text{ pJ}$. During sitting or lying down, input spike density drops by **\$88.4\%**, reducing per-query inference energy from $18.4 \mu\text{J}$ down to **$1.1 \mu\text{J}$** !

5. Temporal Interval Aggregation & Grounded QA Engine

5.1 Symbolic Activity Timeline

Window-level classifications are susceptible to transient classification flicker. I pass predictions through a temporal run-length median filter and merge adjacent matching classifications into an `ActivityTimeline` structure: $\mathcal{K} = \langle a_k, t_{\text{start}}, t_{\text{end}}, \bar{c}_k, f_{\text{dom}}, \sigma^2_{\mathbf{a}}, \sigma^2_{\boldsymbol{\omega}} \rangle$. Short intervals below 4.0 seconds are eliminated as transitional noise.

5.2 Deterministic Symbolic Query Execution (Tasks 1–3)

Rather than relying on Large Language Models to compute numerical arithmetic (which frequently hallucinate durations and counts), Tasks 1, 2, and 3 are handled deterministically: - **Identification:** Returns dominant activity over interval with exact label. - **Verification:** Evaluates binary presence of class $\mathcal{K}(a)$. - **Duration Summation:** Sums interval durations: $T_{\text{total}} = \sum_k (t_{\text{end}}, k) - t_{\text{start}}, k)$. - **Event Counting:** Counts disconnected intervals satisfying $t_{\text{start}}, k+1 - t_{\text{end}}, k > \Delta t_{\text{pause}}$. - **Onset Localization:** Identifies $t_{\text{onset}} = \min_k t_{\text{start}}, k)$.

5.3 Biomechanical Grounded Explanations

Explanations are synthesized from quantified signal physics: - **Walking:** Cites step frequency (1.6–2.0 Hz) and harmonic vertical acceleration (Acc-Z). - **Running:** Cites elevated cadence ($> 2.8 \text{ Hz}$), sharp heel-strike jerk peaks, and wide-amplitude gyroscope arm swing. - **Sitting vs Lying Down:** Differentiates posture by analyzing the static gravity vector ($a_z \approx 9.8 \text{ m/s}^2$ for sitting upright vs $a_x \approx 9.8 \text{ m/s}^2$ for horizontal recumbency).

5.4 Kinematic Open-World Reasoning (Task 4)


For activities outside the 7 supervised labels: - **Wheeled / Pedal-based Locomotion (Bicycling):** Characterized by periodic pedal cadence ($1.2 \text{--} 1.8 \text{ Hz}$) and continuous gyroscopic roll/yaw adjustments without the impulsive vertical heel-strike impacts characteristic of walking or

running ($\max \|\mathbf{a}\| - \bar{a} < 3.5 \text{ m/s}^2$). - **Prolonged Lying Down:** Flagged when continuous stationary quiescence ($\text{Var}(\mathbf{a}) < 0.05 \text{ m}^2/\text{s}^4$, $\text{Var}(\boldsymbol{\omega}) < 0.01 \text{ rad}^2/\text{s}^2$) extends beyond 300 seconds.

6. Experimental Results & Required Figures

6.1 Accuracy Across Question Types (Figure 1)


Overall performance was evaluated over a diverse test bank of 500 questions covering all four tiers.



- **Identification:** 93.4% exact label match.
- **Verification:** 94.8% positive class F1.
- **Duration:** 91.2% within $\pm 10\%$ relative error tolerance.
- **Count:** 88.5% within ± 1 occurrence count.
- **Comparison:** 94.0% correct comparative duration classification.
- **Evidence Grounding:** 89.6% accepted at $\text{IoU} \geq 0.5$.
- **Open-World Reasoning:** 87.5% rubric score ($\geq 4.0 / 5.0$).
- **Overall Macro-QA Accuracy:** 91.3%.

6.2 7-Class Activity Recognition Confusion Matrix (Figure 2)


Figure 2 displays the confusion matrix across all 7 target classes evaluated on windowed sensor streams.



- **Macro-Precision:** 92.4%
- **Macro-Recall:** 91.8%
- **Macro-F1:** 92.1%
- Primary confusions occur between *sitting* and *standing in place* (due to identical near-zero dynamic jerk, separable only by tilt angle) and transitions between *walking* and *running*.

6.3 Accuracy vs Strictness Curve (Figure 3)


The Intersection-over-Union (IoU) acceptance curve traces grounding robustness across thresholds from 0.1 to 0.9.



At the standard benchmark threshold of $\text{IoU} = 0.5$, my grounding module accepts **86.0%** of cited intervals, maintaining over **68%** acceptance even at strict $\text{IoU} = 0.7$, demonstrating tight temporal bounds.

6.4 Accuracy vs Overhead: Pareto Frontier (Figure 4)

To demonstrate edge deployment viability (Extra Credit), I benchmarked four operating configurations on single-query inference:




Architecture	Model Size	CPU Latency	Peak RAM	Energy / Query	Overall QA Acc	Pareto Optimal?
Full 1D-CNN (FP32)	0.85 MB	4.2 ms	18.2 MB	18.4 μ J	92.8%	Yes (High Acc)
Pruned 1D-CNN (30%)	0.62 MB	3.5 ms	16.5 MB	12.8 μ J	91.9%	Dominated
INT8 Quantized 1D-CNN	0.23 MB	1.8 ms	9.4 MB	4.9 μ J	91.5%	Yes (Edge Balanced)
Neuromorphic LIF SNN	0.18 MB	1.2 ms	6.1 MB	1.1 μJ	90.7%	Yes (Ultra-Low Power)

The Neuromorphic SNN reduces energy by **\$16.7\times\$** while incurring only a negligible 2.1% decrease in accuracy, forming the dominant Pareto boundary for wearable microcontrollers.

6.5 Robustness Under Sensor Degradation (Figure 5)

Figure 5 evaluates system resilience when sensor packets are randomly dropped (0% to 50% packet loss).



With Butterworth filtering and piecewise linear interpolation, system QA accuracy remains above **87.8%** even at 30% sample loss, whereas an un-interpolated baseline collapses below 62%.

7. Canonical Grounded Case Studies

Case Study 1: Task 2 Duration Summation

Query: "How long was the user walking?"

Answer: 700 seconds

Activity/Event: Walking

Evidence:

Timestamp(s): 905 to 1420, 2110 to 2295 (seconds from start)

Sensor Modality: Accelerometer, Gyroscope

Sensor Channel(s): All

Explanation: Walking was detected in 2 separate intervals, of 515 and 185 seconds, which

Case Study 2: Task 3 Evidence Grounding

Query: "Did the user begin running at any point, and if so, when?"

Answer: Yes, running began at 1512 seconds

Activity/Event: Onset of running

Evidence:

Timestamp(s): 1512 to 1980 (seconds from start)

Sensor Modality: Accelerometer, Gyroscope

Sensor Channel(s): All

Explanation: A sustained rise in accelerometer magnitude at a higher step frequency, to

Case Study 3: Task 4 Open-World Locomotion

Query: "Was the user using a wheeled or pedal-based mode of movement?"

Answer: Yes

Activity/Event: Unknown outdoor physical activity, consistent with cycling

Evidence:

Timestamp(s): 2400 to 3120 (seconds from start)

Sensor Modality: Accelerometer, Gyroscope

Sensor Channel(s): All

Explanation: The segment shows smooth, continuous, cyclic acceleration at a steady cadence

8. Reproducibility & Deployment Instructions

1. **Clone and Install:** `bash git clone <repository_url> cd ask_the_sensors pip install -r requirements.txt`

2. **Execute On Fresh Recording:** `bash python run_qa.py --data data/sample_recording_25hz.csv --query "How long was the user walking?"`
3. **Recreate Figures:** `bash python generate_figures.py`
-

9. Individual Author Contribution Statement

This project was independently conceived, designed, engineered, and evaluated by the sole author, **Anubhav**:

- **System Architecture & Ideation:** Formulated the multi-tier sensor question-answering framework and originated the novel integration of an event-driven Neuromorphic Spiking Neural Network (SNN) with delta-spike modulation to overcome the battery drain of continuous IMU monitoring during sedentary periods.
- **Signal Preprocessing & Physics Featurization:** Formulated the uniform 25 Hz resampling and interpolation mathematics, Butterworth artifact filtering, and the mathematical extraction of 20 biomechanical features (FFT cadence power spectra, dynamic jerk variance, Signal Magnitude Area, and gravitational pitch/roll tilt).
- **Neural & Neuromorphic Modeling:** Architected the dual-backbone system: the baseline 1D-CNN, its dynamic INT8 quantized edge variant, and the custom PyTorch Leaky Integrate-and-Fire (LIF) network featuring FastSigmoid surrogate gradient backpropagation and Synaptic Operation (SynOps) energy modeling.
- **Symbolic Reasoning & Open-World Kinematics:** Designed the temporal interval aggregation logic (run-length median filtering) to eliminate LLM arithmetic hallucinations, engineered the deterministic query answering algebra for Tasks 1–3, and formulated the kinematic body-dynamics rules for Task 4 (cycling pedal cadence vs. running heel strikes; prolonged recumbency).
- **Evaluation Suite & Reporting:** Authored the benchmark evaluation harness, generated all 5 mandatory high-resolution figures, and authored this comprehensive technical report.

10. AI-Use Disclosure & Integrity Statement

In accordance with the CS60055 academic integrity policies announced for Hackathon Challenge 1, this statement provides full transparency regarding the role of AI tools during the project lifecycle:

- **Originality of Ideation and Technical Formulation:** All problem framing, architectural design decisions, mathematical models (LIF membrane differential dynamics, synaptic energy scaling, and temporal interval overlap logic), heuristic thresholds, and investigative hypotheses were conceived, formulated, and directed solely by the author.
- **Academic Sources and Literature Foundation:** The technical design draws upon and synthesizes established peer-reviewed literature in ubiquitous computing and neuromorphic signal processing:
- *Vaizman et al. (IEEE Pervasive Computing, 2017)*: Provided the ExtraSensory dataset context, sampling characteristics, and multi-label behavioral taxonomies.
- *Maass (Neural Networks, 1997)* and *Neftci et al. (IEEE Signal Processing Magazine, 2019)*: Informed the mathematical formulation of Leaky Integrate-and-Fire spiking neurons, delta-

modulation event thresholds, and surrogate gradient backpropagation for neuromorphic edge efficiency.

- *James F. Allen (Communications of the ACM, 1983)*: Provided the temporal interval logic underpinning the deterministic query execution over activity timelines to avoid generative hallucinations.
- *Biomechanical Gait Dynamics Literature*: Informed the frequency-domain cadence bands (1.5–2.0 Hz walking cadence, >2.8 Hz running cadence) and the harmonic distinction between continuous circular pedaling in cycling versus impulsive heel strikes in running.
- **Specific Role of AI Assistance**: An AI assistant was employed strictly as an interactive coding accelerator, akin to an advanced compiler assistant and documentation formatter:
- Drafting standard Python API boilerplate (e.g., standard parameter calls for `scipy.signal.butter` and Matplotlib plot styling).
- Assisting with Markdown syntax layout and formatting.
- **Verification and Plagiarism Assurance**: No code, data, or experimental results were copied or reused from unauthorized external sources or peer groups. No synthetic data was represented as unverified real-world measurements. All code was locally tested, debugged, and executed end-to-end by the author, and all benchmark figures were programmatically produced from the validated codebase.

References

1. Y. Vaizman, K. Ellis, and G. Lanckriet, "Recognizing Detailed Human Context In-the-Wild from Smartphones and Smartwatches," *IEEE Pervasive Computing*, vol. 16, no. 4, pp. 62–74, 2017.
2. ExtraSensory Dataset Official Repository: <http://extrasensory.ucsd.edu/>
3. W. Maass, "Networks of spiking neurons: the third generation of neural network models," *Neural Networks*, vol. 10, no. 9, pp. 1659–1671, 1997.
4. E. O. Neftci, H. Mostafa, and F. Zenke, "Surrogate Gradient Learning in Spiking Neural Networks," *IEEE Signal Processing Magazine*, vol. 36, no. 6, pp. 51–63, 2019.
5. J. F. Allen, "Maintaining Knowledge about Temporal Intervals," *Communications of the ACM*, vol. 26, no. 11, pp. 832–843, 1983.